


OpenCL Introduction


PRACE/LinkSCEEM Winter School January, 26th 2011



Agenda

- Introduction
- OpenCL Architecture
- OpenCL Application

Introduction



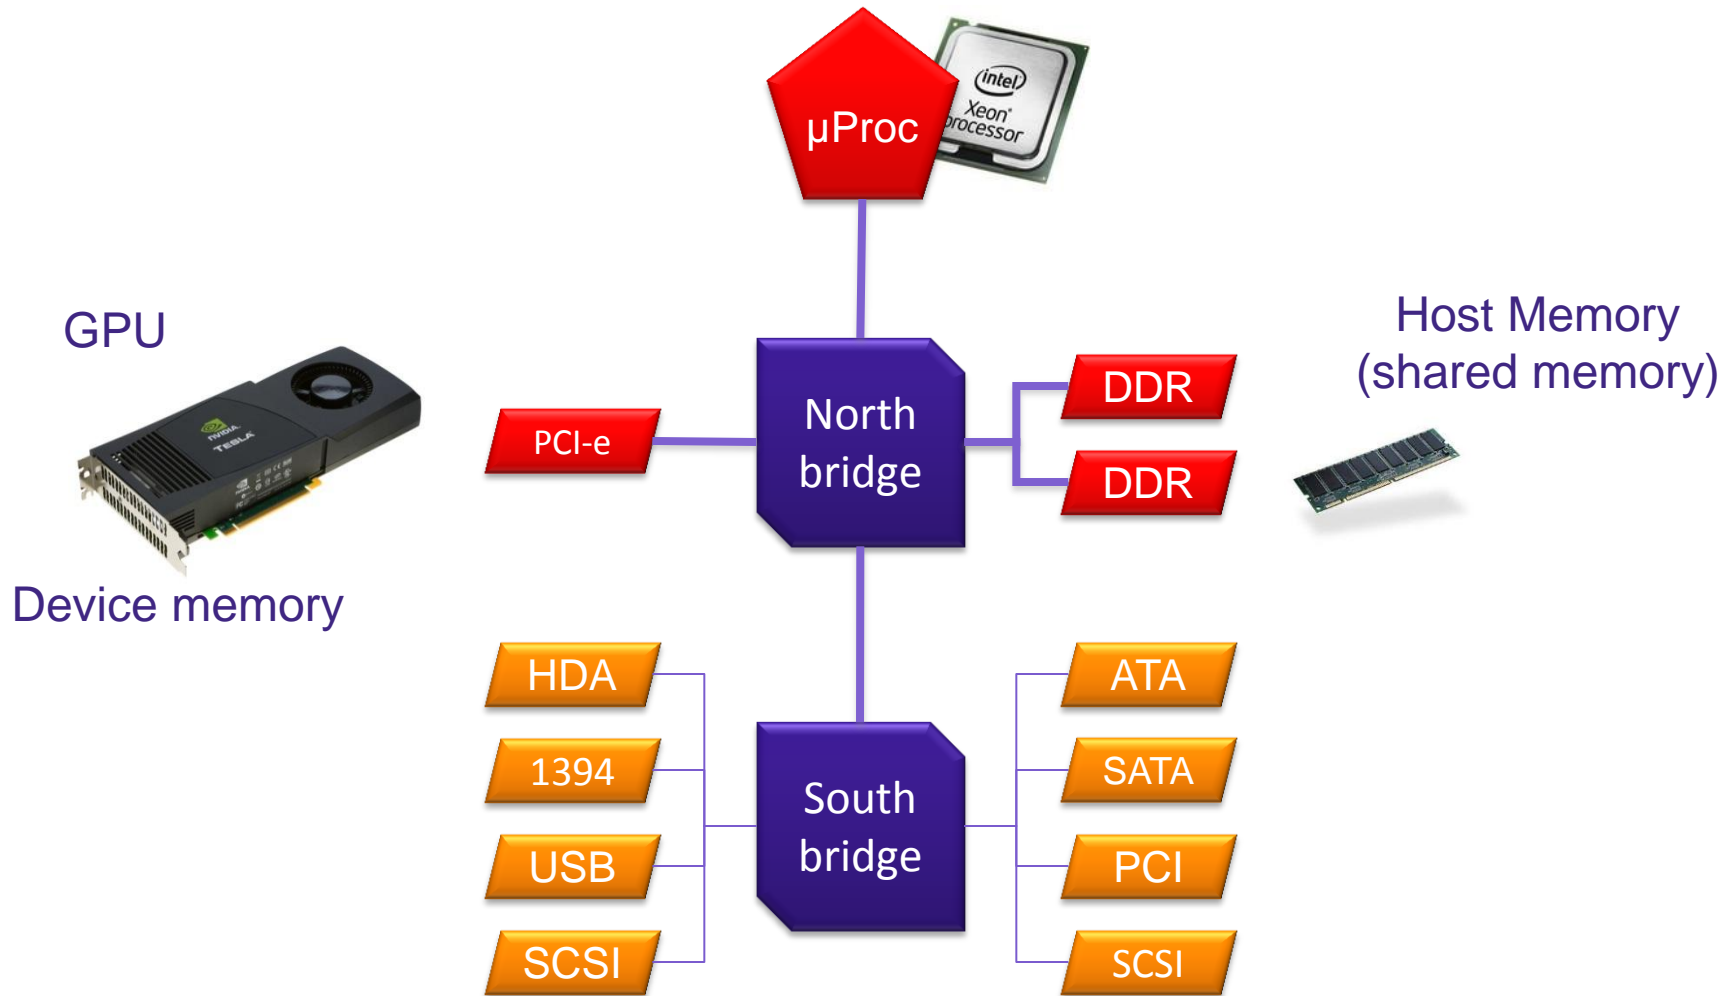
Before OpenCL

- **GPGPU**
 - Vertex / pixel shaders
 - Heavily constrained and not adapted
- **Brook**
 - Then Brook+
 - Then CAL/IL
- **CUDA**
 - Widely broadcasted
- **No one of these technologies is hardware agnostic**
 - Portability is not possible

What is Hybrid Computing with OpenCL?

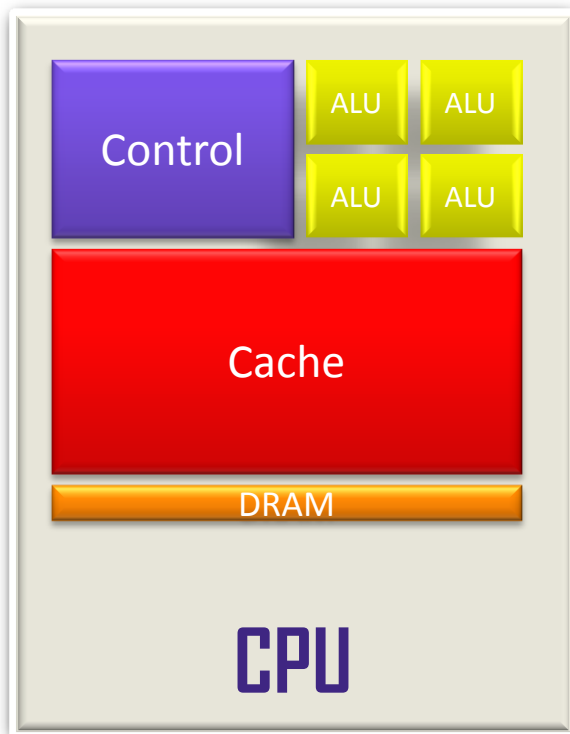
- OpenCL is
 - Open, royalty-free, standard
 - For cross-platform, parallel programming of modern processors
 - An Apple initiative
 - Approved by Intel, Nvidia, AMD, etc.
 - Specified by the Khronos group (same as OpenGL)
- It intends to unify the access to heterogeneous hardware accelerators
 - CPUs (Intel i7, ...)
 - GPUs (Nvidia GTX & Tesla, AMD/ATI 58xx, ...)
- What's the difference with CUDA or CAL/IL?
 - Portability over Nvidia, ATI, S3... platforms + CPUs

Heterogeneous Platforms Architecture




CPU vs GPU Architecture

- Computational kernels would be different



General purpose architecture



Massively data parallel

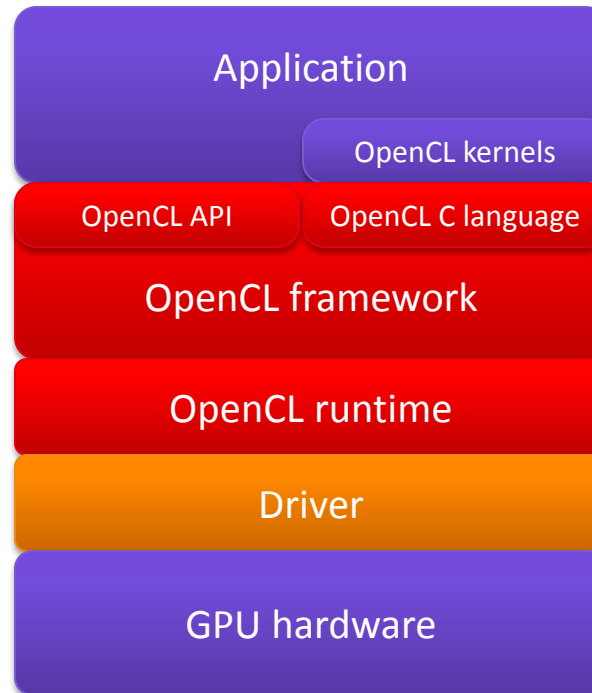
OpenCL Devices

- Intel & AMD
 - X86 w/ \geq SSE 3.x
- S3
 - NV1000
 - 5400E
 - ...
- IBM Cell
- NVIDIA
 - All CUDA cards
 - But not all the drivers
- ATI
 - Radeon & Radeon HD
 - FirePro, FireStream
 - Mobility...



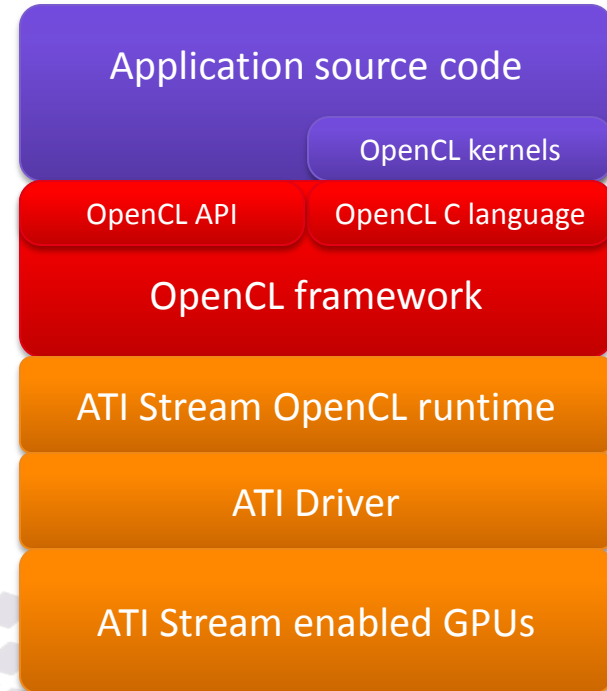
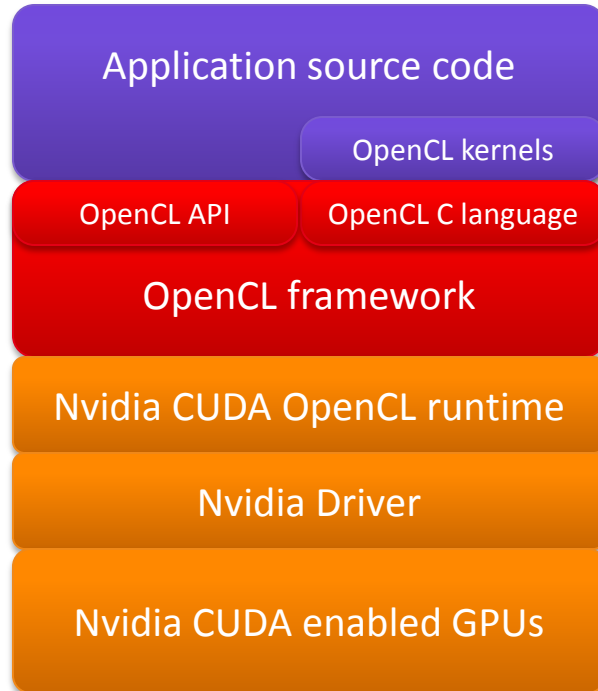
Inputs/Outputs with OpenCL programming

- OpenCL architecture



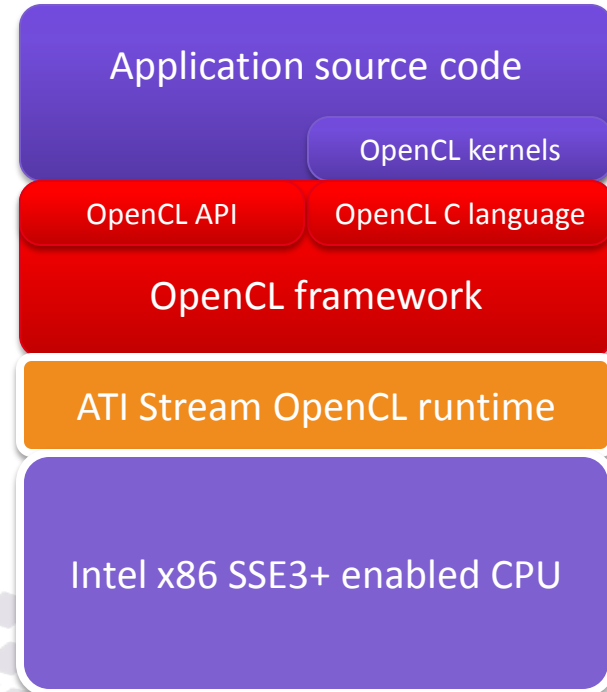
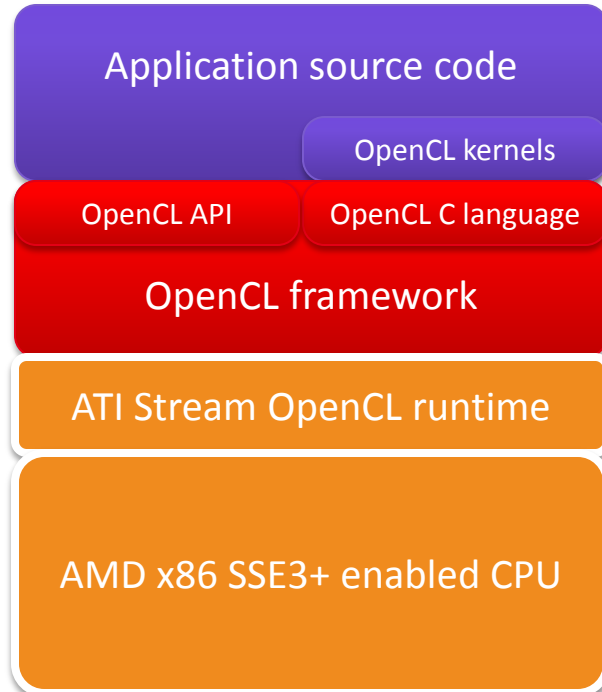
Inputs/Outputs with OpenCL programming

- Nvidia GPUs
- Ati GPUs

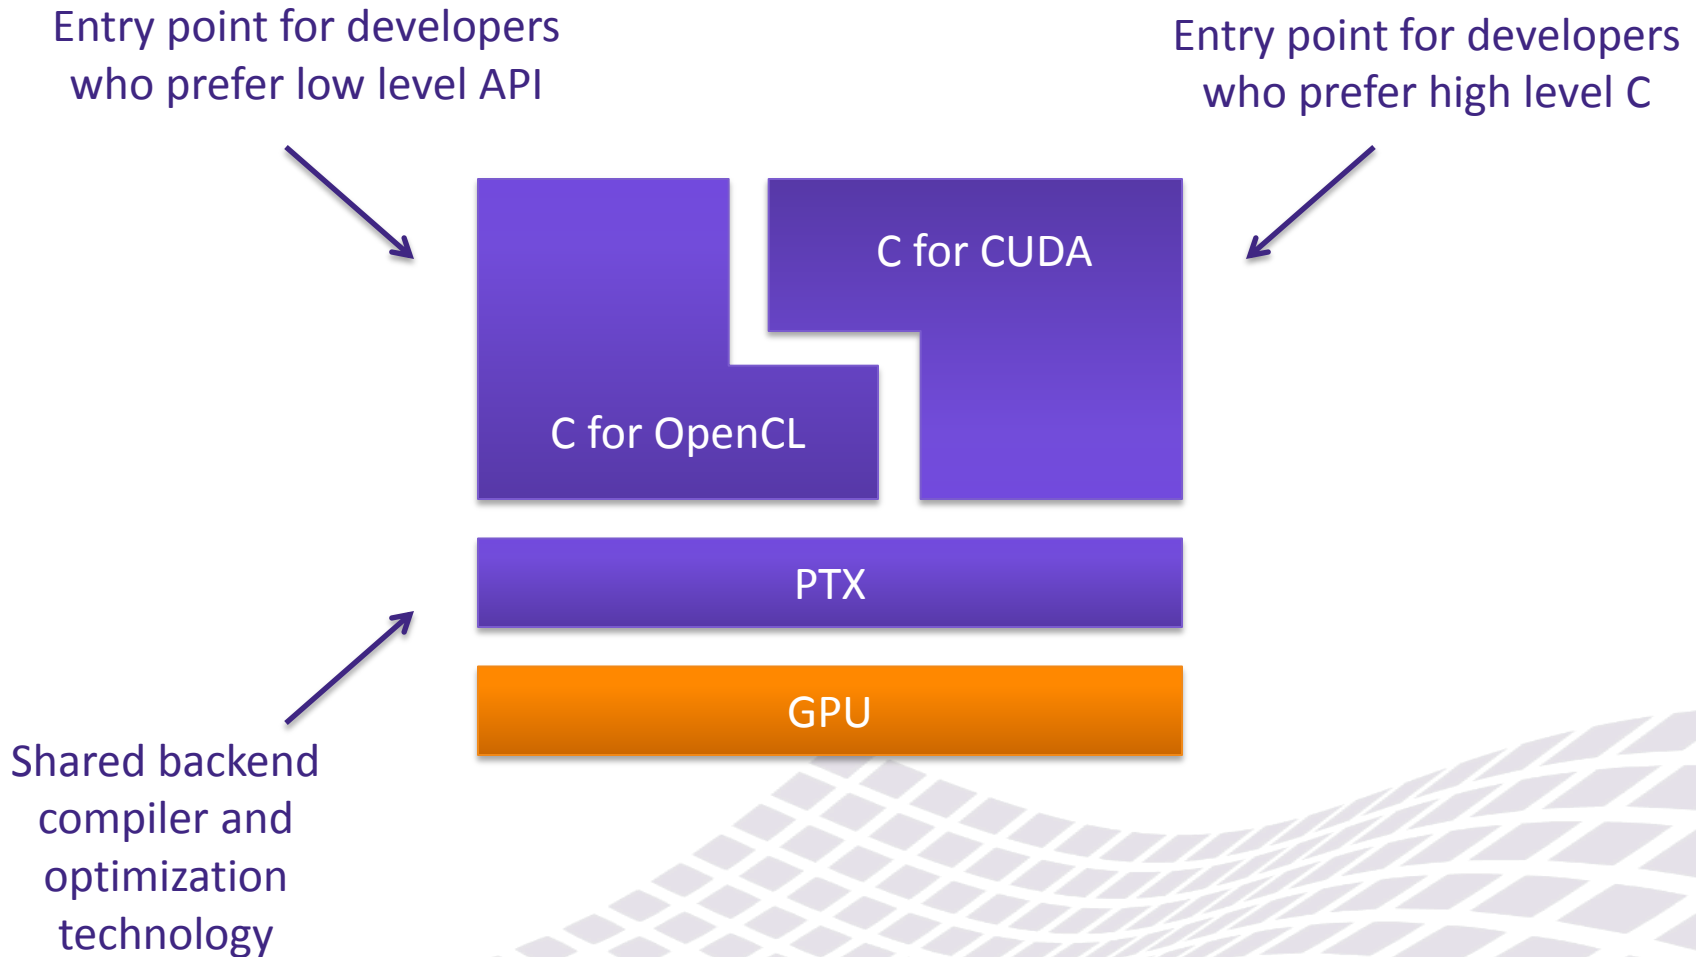


Inputs/Outputs with OpenCL programming

- AMD processors
- Intel processors




OpenCL and C for CUDA




- **Compilation and execution**
 - Include header and link with OpenCL library
 - On ATI, export DISPLAY=:0.0
- **C language API**
 - Bindings C++ (official)
 - Bindings Java
 - Bindings Python
 - ...
- **Extensions exist to**
 - OpenGL
 - Direct3D

OpenCL Architecture



Platform Model

- Model consists of one or more interconnected devices



- Computations occur within the Processing Elements of each device

Platform Version

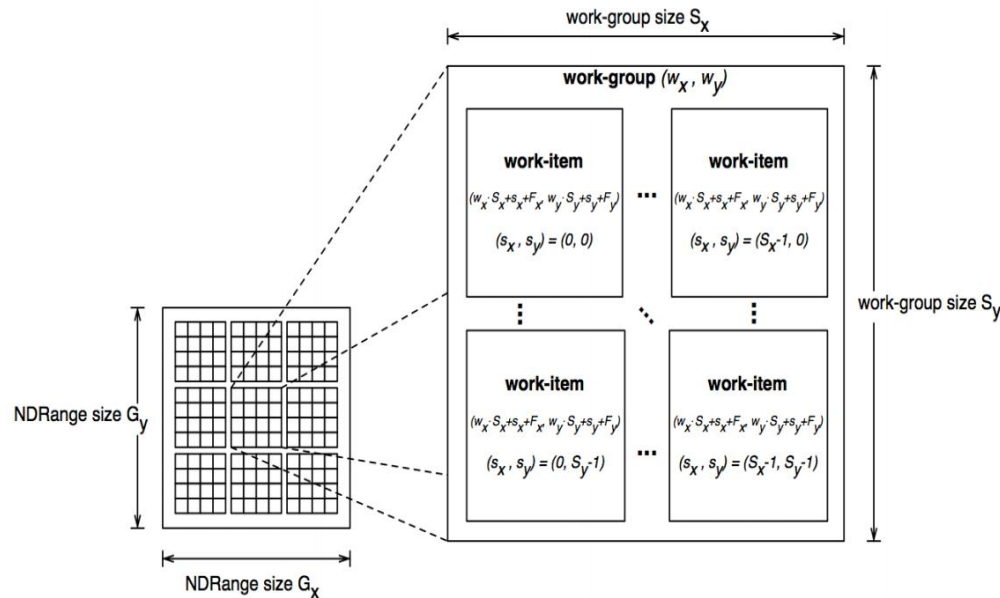
- 3 different kind of versions for an OpenCL device
- The platform version
 - Version of the OpenCL runtime linked with the application
- The device version
 - Version of the hardware (driver)
- The language version
 - Higher revision of the OpenCL standard that this device supports

- Kernels are submitted by the host application to devices throw command queues
- Kernel instances, called Work-Item (WI), are identified by their point in the NDRange index space
 - This enables to parallelize the execution of the kernels
- But still 2 programming models are supported
 - Data-parallel
 - Task parallel
- So even if we have a single programming model, we should have two different programming approaches according to the paradigm we are considering

- **Data-parallel**
 - A frequent way is a one-to-one mapping between the elements of a memory object and the WI space a kernel can implement in parallel
- **Task parallel**
 - Single instance of kernels are executed independly of any NDRange
 - Equivalent to work-group contains only one work-item
- The programmer then defines the number of work-items to execute and their allocation into work-groups
- **2 ways of synchronization**
 - On the hardware: WI of a same WG can synchronize their execution
 - On the host: commands queued up to the Command Queue for a same context can be executed asynchronously, in order or not

- CPU cores can handle only a few tasks
 - But more complex
 - Hard control flows
 - Memory cache
 - Different tasks
 - Or small computation grid (NDRange)
- GPU threads are extremely lightweight
 - Very little creation overhead
 - Simple and regular computations
 - GPU needs 1000s of threads (w.i.) for full efficiency

- NDRange is a N-Dimensional index space
 - N is 1, 2 or 3
 - NDRange is defined by an integer array of length N specifying the extent of the index space on each dimension



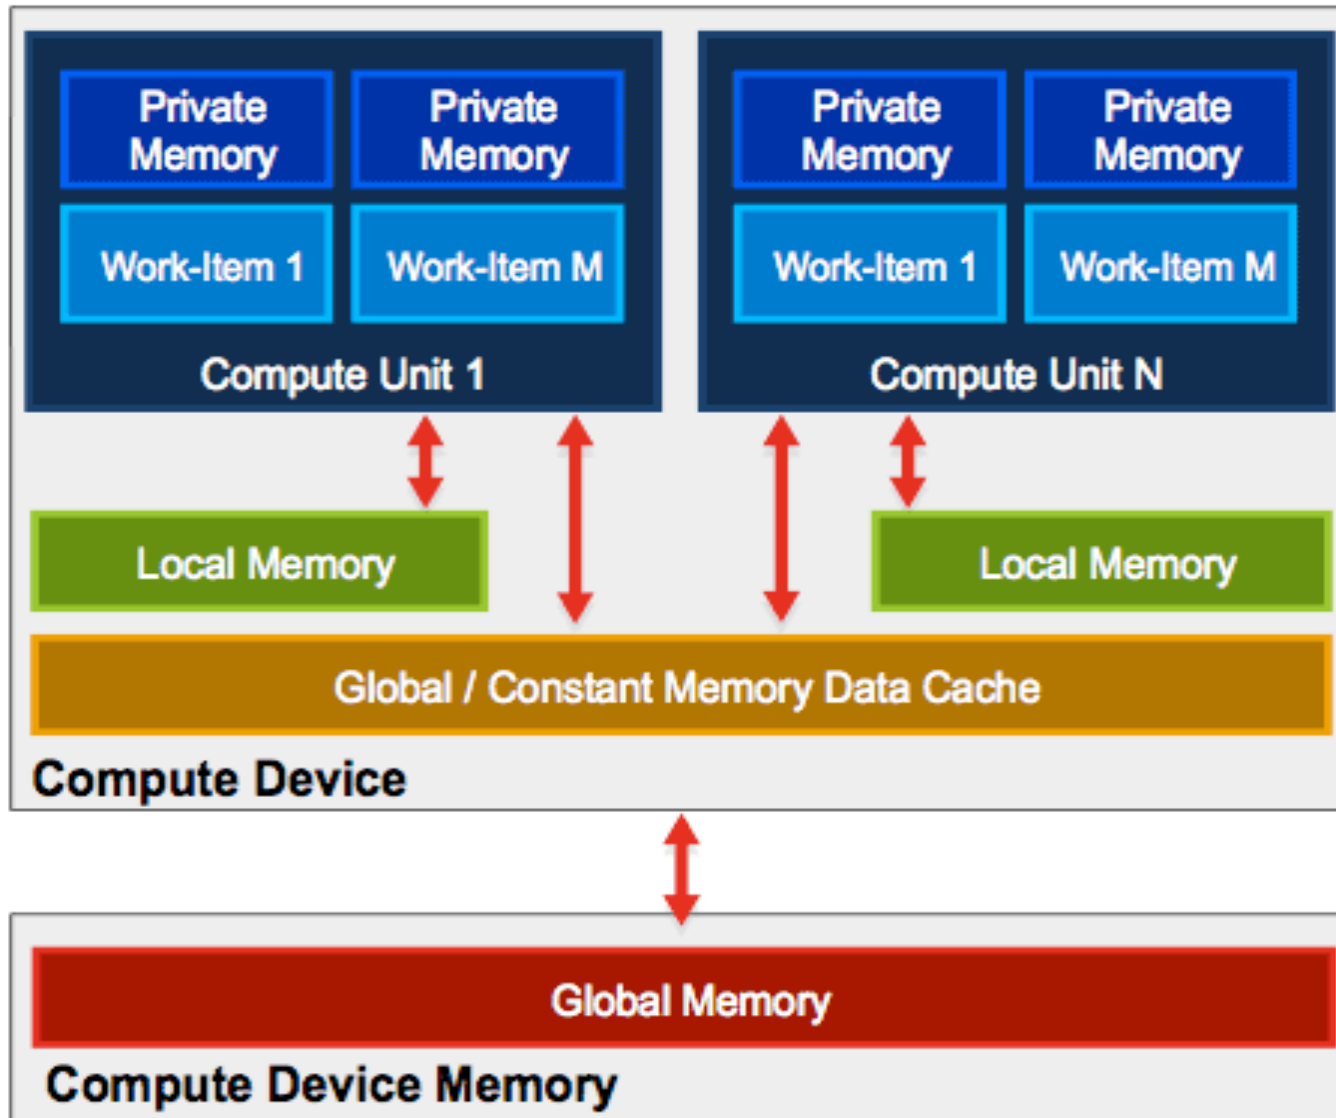
Work-Groups & Work-Items

- Work-Items are organized into Work-Groups (WG)
- Each Work-group has a unique global ID in the NDRange
- Each Work-item has
 - A unique global ID in the NDRange
 - A unique local ID in his work-group

Memory Model

- Four distinct memory regions
 - Global Memory
 - Local Memory
 - Constant Memory
 - Private Memory
- Global and Constant memories are common to all WI
 - May be cached depending on the hardware capabilities
- Local memory is shared by all WI of a WG
- Private memory is private to each WI

Memory Architecture



Memory Allocation

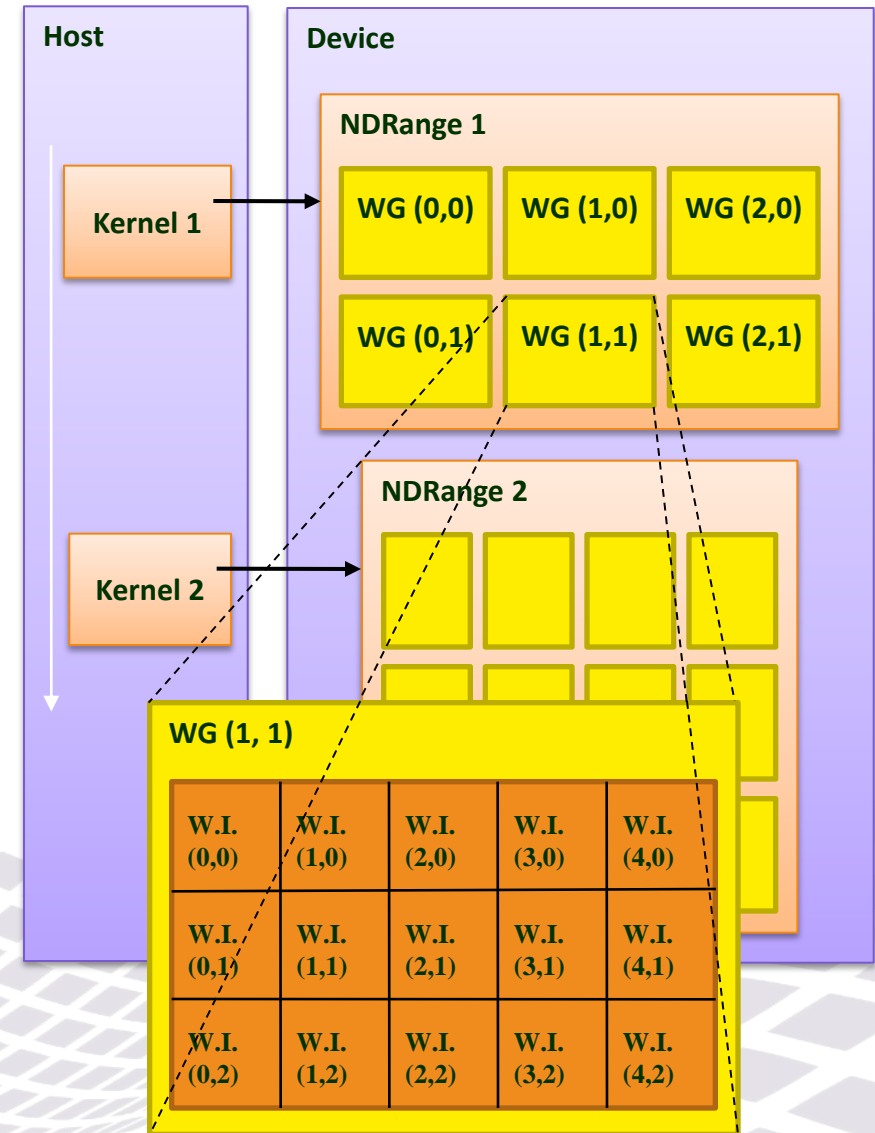
- Allocations depends on user's point of view

	Global	Constant	Local	Private
Host	Dynamic Allocation	Dynamic Allocation	Dynamic Allocation	No Allocation
	R/W	R/W	No Access	No Access
Kernel	No Allocation	Static Allocation	Static Allocation	Static Allocation
	R/W	Read-only	R/W	R/W

- **Qualifiers**
 - `__global` for global memory
 - `__local` for local memory
 - `__constant` for constant memory
 - `__private` for private memory
- The "`__`" prefix is not required before the qualifiers
- **Inside a kernel**
 - By default variables are in private memory
 - By default, variables declared as pointer points to private memory
- Pointers passed as arguments to a kernel function must be of type `__global`, `__constant`, or `__local`

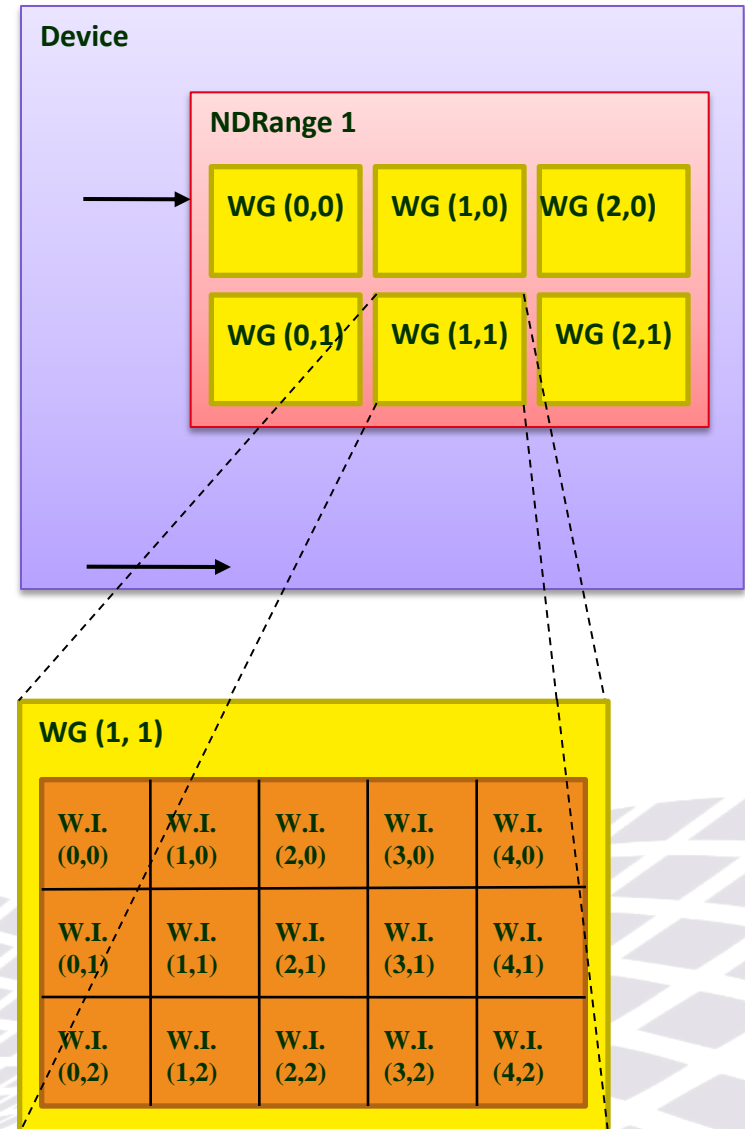
Thread Batching: Work Groups and NDRange

- A kernel is executed as a grid of work groups (NDRange)
 - All W.I. can share a data memory space inside their work group
- A work group is a batch of threads which can cooperate with each other by:
 - Synchronizing their execution
 - For hazard-free local memory accesses
 - Efficiently sharing data through a low latency local memory
- Two threads from two different WG cannot cooperate




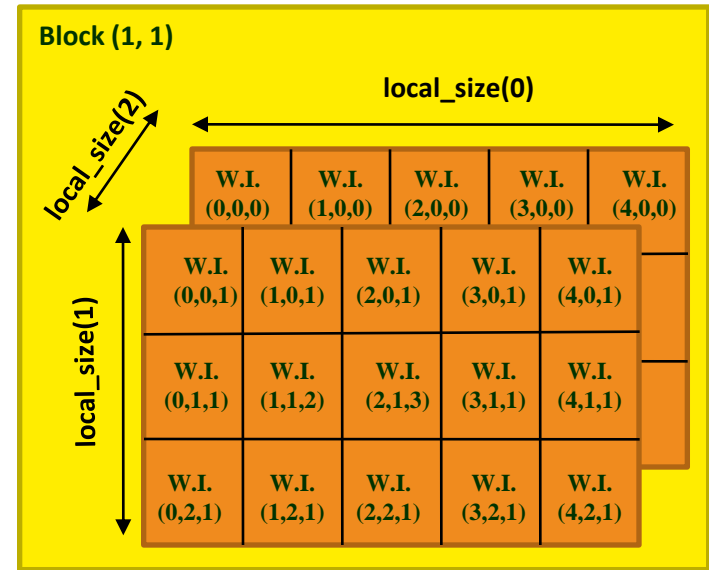
Work groups and Work Item IDs

- Work items and work groups have IDs
 - So each thread can decide what data to work on
 - W.G. ID: ND
 - W.I. ID: ND
- Simplifies memory addressing when processing multidimensional data
 - Image processing
 - Solving PDEs on volumes
 - ...




Work Group and Work Item built-in functions

- Work Item keywords
 - `get_local_id(dim)` returns the index inside the W.G.
 - `get_global_id(dim)` returns the index inside the NDRange
 - `get_local_size(dim)` returns the block size
 - `get_global_size(dim)` returns the NDRange size
- Work Group keywords
 - `get_group_id(dim)` returns the work group index
 - `get_num_groups(dim)` returns the grid size



OpenCL Application



- OpenCL platform = host + OpenCL devices
 - Devices can be CPUs, GPUs and so...
 - At the moment, Nvidia SDK does not implement OpenCL on CPUs
 - ATI allows to use OpenCL either on GPUs or x86 CPUs (what are AMD and Intel's)
- Available OpenCL devices depend on their drivers
 - Only the driver is going to tell OpenCL of the capabilities of its hardware
 - CPUs have no driver
- Drivers for Nvidia and ATI GPUs cannot be installed on a same host
 - So this is one platform or another
 - But take care, maybe one day it will be possible

The execution context

- Get the list of all platforms
 - Return the list of OpenCL driver installed on the system
 - `clGetPlatformIDs`
 - `clGetPlatformInfo`
- Get the list of devices from a platform
 - Get list of available devices of a platform
 - `clGetDeviceIDs`
 - `clGetDeviceInfo`
- An execution context
 - Depends on a list of devices
 - `clCreateContext`

OpenCL Programs

- An OpenCL program consists of a set of kernels
- A program object contains
 - An associated context
 - A program source code
 - The kernel objects associated to
- Two ways to create a OpenCL program
 - From a string containing the kernel code
 - From a binary file (specific to an architecture)
- The “__kernel” keyword identifies a kernel

Creating Program & Kernel Objects

- The source program is a string
 - Contains kernels
 - Reads a file at runtime
 - Allows to modify kernels without recompiling the application
 - `clCreateProgramWithSource`
- Build the program object
 - Compile the program for a list of devices
 - `clBuildProgram`
- How to get kernels objects ?
 - Depends of program object
 - `clCreateKernel`

Kernel Syntax

- In the kernel, the language must be OpenCL C
- Built-in Functions
 - `int get_local_id(int dim)`
 - `int get_global_id(int dim)`
 - `int get_local_size(int dim)`
 - `int get_global_size(int dim)`
 - ...
- Math functions
 - Trigonometric functions
 - Mad function
 - Min and max functions
 - Exponential, logarithm and power
 - modulus
 - ...

The Command Queue

- An OpenCL command-queue
 - Used to queue a set of operations
 - Associated to a unique device
 - In-order or out-of-order
 - `clCreateCommandQueue`
- To synchronize a command queue
 - `clFlush`
 - `clFinish`
- Having multiple command-queues allows applications to queue multiple independent commands without requiring synchronization


Device Memory

- Two types of memory objects
 - Buffer Objects : 1D memory
 - Image Objects : 2D-3D memory
- Any data transfer to/from the device implies
 - A host pointer
 - A device memory object
 - The size of data (in bytes)
 - The command queue to enqueue
 - If it is a blocking transfer or not
- To transfer to/from the device
 - `clEnqueueWriteBuffer`
 - `clEnqueueReadBuffer`

- Set kernel arguments
 - Cause of kernel compilation at runtime
 - `clSetKernelArg`
- 2 way to launch kernels
 - Enqueue task kernel
 - Enqueue NDRange kernel
- Set kernel options
 - Set global size of NDRange
 - Set local size of workgroups
- Kernel launches are asynchronous

OpenCL Kernel Performance


PRACE/LinkSCEEM Winter School January, 26th 2011



Agenda

- OpenCL Image Objects
- Local Memory
- Transformations & Optimizations
- Non-portability of the performances

OpenCL Image Objects



- Allows to allocate 2D or 3D arrays in a specific object
 - Allow to exploit some hardware functions
 - Based on the Textures in OpenGL
- Allows vector types, by using specific image format
 - The most common is float4, but float, float2 and various integer types are possible
 - List of available formats is implementation-dependant
 - 3D images are not required for Embedded Profile
- Elements of an image are stored in an opaque format
 - Cannot be directly accessed using a pointer

- Memory object argument of a kernel needs a qualifier
 - `__read_only` (by default)
 - `__write_only`
 - Kernel cannot read from and write to the same image memory object
- Like with OpenGL textures, many options for the sampler
 - Normalized coordinates or not
 - Address can be clamped or in tiling mode
 - Results can be the nearest value, or a linear interpolation
- One of the many, many functions:

```
float4 read_imagef (image2d_t image,  
                    sampler_t sampler,  
                    int2 coord);
```


```
void write_imagef (image2d_t image,  
                  int2 coord,  
                  float4 color);
```

To use or not to use Images

- Images cannot be updated (read & write in the same kernel)
 - useless for instance for the “C” matrix in xGEMMs
- Therefore, they can be cached very effectively when reading : no cache coherency needed
 - In particular, images are very, very good on ATI 58xx
 - Not so much on Nvidia C2050, as global memory is also cached
- Images are not necessarily native to all hardware
 - Neither IBM Power 7 nor IBM Cell B/E support Images
- Can be extremely efficient if the linear interpolator can be leveraged
 - Unfortunately not a very common thing in HPC codes

- OpenCL specifications part 6.1.2 (types), 6.4 (operations)
- Vector data types (float4, ...) represent a pseudo-array of elements (“components” in OpenCL)
- OpenCL specifies per-component semantic
 - So the product of two float4 is another float4, not a scalar float : it is not a dot product
 - Scalar are implicitly converted into the proper type by replicating the value
- Much more efficient on vectorized architecture, as only $\frac{1}{4}$ (for float) of the peak is reachable in scalar mode
- Component can be accessed in different ways
 - Scalar: temp.s0, temp.s1
 - Sub-vector: temp.odd, temp.left

Local Memory




- OpenCL specifications, part 3.3
- Local memory is local to a work-group
 - Shared by all work-items of work-group
 - Can be used for sharing between work-items
 - Can be used to cache global memory
- Low latency access
 - On Nvidia C1060-class hardware, “Nvidia shared memory” has the same latency as registers
- 2 ways to allocate it
 - Statically, inside the kernel
 - Dynamically, from the host as a parameter

Local Memory Usage

- Many work-items of a work-group have access to the same subset of data
 - Allows spatial locality between work-items
 - Allows temporal locality in one or more work-items
 - On Cell, is the only way to efficiently access data
- Can be used to realign memory
 - On Nvidia, avoid uncoalesced access in global memory by pre-loading aligned block into local memory
- To synchronize work-items of a work-group
 - Optional atomic operations in local memory (part 9.6)
- Be careful about bank conflict on Nvidia Hardware

- Inside a kernel
 - Declare local memory as a static array
 - Use the keyword as a type prefix
 - `__local` or `local`
- Declare a Kernel with parameter in local memory
 - Allocation done during set of kernel arguments

Transformations & Optimizations



Memory Transfer Optimization

- 2 types of transfers
 - Blocking (“synchronous”)
 - Non-Blocking (“asynchronous”)
- In the function `clEnqueueRead/Write`, set the parameter **blocking**
 - `CL_TRUE`, make a blocking transfer
 - `CL_FALSE`, make a non-blocking transfer
- For a non-blocking transfer
 - Needed to control asynchronism
 - Allows to synchronize commands, command queues and the host

- On Nvidia GPU
 - Global memory access > 400-700 cycles
 - Prefers coalesced, small-size accesses (4 bytes per threads)
- On AMD/ATI GPU
 - Data reuse from Global (readings cached/writings not cached)
 - Images can be very fast
 - Prefers access to global in larger chunks (16 bytes, i.e. float4 per threads)
- Optimize the use of register
 - Depends of the number of blocks on Nvidia Hardware
- Use the local memory
 - Low latency access

NDRange Optimization

- Optimization of the NDRange is hardware dependant
- Required for efficient use of GPUs
 - Must combine the global & local size for maximum efficiency of the various cores & the memory subsystem
- On Nvidia GPU, WG size must be multiple of warp size
 - Like with CUDA try different size
 - 16x16
 - 32x4
 - 64x1
 - ...
- If it works well in CUDA, then it should also works well in OpenCL :-)
 - Each generation of hardware has different requirements

- On ATI GPU, WG size must be multiple of wavefront size
 - More WI can be generated
 - GPU can switch when a WI is stalled
 - Hiding memory latency
 - Optimizing execution
- On IBM Cell, the number of WG must be a multiple of the number of SPU
 - work-item of a work-group are merged (unrolled) in a single thread, executing on a single SPU
 - Large multiple are not necessarily better, having the exact number can be the most efficient choice

- **Loop Unrolling**
 - Beware of rest loop
 - Fullunroll for small loops
 - Split space iteration (Nvidia Hardwares)
- **Software pipelining**
 - Increases the distance between a memory load and its use
 - hides part of the latency
- **Avoid conditional**
 - On GPU, avoid divergent branches
 - Use predication rather than control-flow when possible

Native functions


- Native functions are implementation defined
 - native_* prefix identifier
- Native functions are faster than built-in functions
 - But precision is implementation-dependant
- Available operation
 - Division
 - Arithmetic : cos, sin, tan
 - Logarithm : base-e, 2, 10
 - Exponential : base-e, 2, 10
 - Square root and inverse

- Global memory optimization
 - Think about coalescing : alignment, granularity, size of accesses
 - Access to shared memory without bank conflicts
 - Ideal pattern has changed on Fermi, with 32 banks instead of 16
- NDRange optimization
 - Like CUDA grid
- Use native_* math functions where possible
- Partition the computation
 - nbWorkgroups >> nbMultiprocessor
 - Keep resources low enough ; the register file & the local memory are shared by all work-item in all work-groups executing simultaneously on a multiprocessor

Hardware dependant : ATI/AMD GPU

- Think about vectorization
 - float4 is best, in particular for global memory accesses
 - Helps in computation also
- NDRange optimization
 - 64 work-item per work-group
 - Multiple work-item per wavefronts
- Data reuse
 - Reads are cached
 - But not writes

Non-portability of the performances



The case of SGEMM


- Classic, well-known operation
 - Should be extended to DGEMM when double-precision support becomes available on all platforms...
- Quick study in OpenCL, comparing on ATI GPU, Nvidia GPU, IBM Cell B/E
 - The CPU is the more exotic of the three :-)
- Completely different code in each case !
- Kernel validated on GPUs by integration in HMPP codelet

- Kernel from the GATLAS <http://golem5.org/gatlas/> OpenCL generator
 - Generates many forms of vectorized OpenCL SGEMM kernels, adapted to each problem size
- Non-vectorized kernel are very inefficient on ATI hardware
- 8x8 is the default, and the most efficient work-group size
- Grid size is large, and depend on internal unrolling in the generator
- Using images for input matrices A & B, is much, much better
- On $n=m=k=4000$ (4096 is noticeably less efficient)
 - About 500 Gflops with buffer on HD5870
 - About 1.3 Tflops with images !

- Using GATLAS-derived kernels (current version of GATLAS doesn't support non-vectorized kernels, work is in progress) or HMPP-generated kernels
- vectorized kernels are inefficient on C1060
 - Float4 memory accesses don't coalesce very well, and generates lots of local memory bank conflicts
 - Execution units are not vectorized
- Non-vectorized kernels are good
 - Opposite reason !
- Compared to ATI, larger unrolling factor
 - Compensating for lack of vectorizations
- Images couldn't be tested

- Best non-OpenCL matmult on Cell at <http://www.tu-dresden.de/zih/cell/matmul>
- Achieve >99% of peak under ideal conditions...
 - Assembly code, optimized DMA transfers, double buffering, ...
- Current OpenCL Kernel only 65% of peak :-(
- Requires use of `async_work_group_copy` (not used on GPU) to use DMA for global -> local copies
- Requires massively unrolled float4 arithmetic
- Also partial double buffering, software pipelining to hide DMA latencies, ...
- Much more similar to traditional CPU optimizations

Hands-on : Convolution



- 6 incremental steps :
 - Communication part in function.cc
 - Naive Kernel
 - Kernel with static Local memory
 - Kernel with dynamic Local memory
 - Full unroll loops
 - Put hard coded coefficient
- Convolution :

$$\begin{aligned} B[i][j] = & 2 * Cst0 * A[i][j] \\ & + Cst1 * (A[i-1][j] + A[i+1][j] + A[i][j-1] + A[i][j+1]) \\ & + Cst2 * (A[i-2][j] + A[i+2][j] + A[i][j-2] + A[i][j+2]) \\ & + Cst3 * (A[i-3][j] + A[i+3][j] + A[i][j-3] + A[i][j+3]) \\ & + Cst4 * (A[i-4][j] + A[i+4][j] + A[i][j-4] + A[i][j+4]) \end{aligned}$$

Vasnier Jean-charles
Application Engineer
jvasnier@caps-entreprise.com
+33 (0) 222 511 600

